

## 3.08.2023(P2)

```
In [79]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [80]: df=pd.read_csv(r"C:\Users\user\Downloads\csvs_per_year\csvs_per_year\madrid_2006\madrid_2006.csv")
df
```

Out[80]:

|        | date                | BEN  | CO   | EBE  | MXV   | NMHC | NO_2       | NOx        | OXY  | O_3   | PM10      |
|--------|---------------------|------|------|------|-------|------|------------|------------|------|-------|-----------|
| 0      | 2002-04-01 01:00:00 | NaN  | 1.39 | NaN  | NaN   | NaN  | 145.100006 | 352.100006 | NaN  | 6.54  | 41.990002 |
| 1      | 2002-04-01 01:00:00 | 1.93 | 0.71 | 2.33 | 6.20  | 0.15 | 98.150002  | 153.399994 | 2.67 | 6.85  | 20.980000 |
| 2      | 2002-04-01 01:00:00 | NaN  | 0.80 | NaN  | NaN   | NaN  | 103.699997 | 134.000000 | NaN  | 13.01 | 28.440001 |
| 3      | 2002-04-01 01:00:00 | NaN  | 1.61 | NaN  | NaN   | NaN  | 97.599998  | 268.000000 | NaN  | 5.12  | 42.180000 |
| 4      | 2002-04-01 01:00:00 | NaN  | 1.90 | NaN  | NaN   | NaN  | 92.089996  | 237.199997 | NaN  | 7.28  | 76.330002 |
| ...    | ...                 | ...  | ...  | ...  | ...   | ...  | ...        | ...        | ...  | ...   | ...       |
| 217291 | 2002-11-01 00:00:00 | 4.16 | 1.14 | NaN  | NaN   | NaN  | 81.080002  | 265.700012 | NaN  | 7.21  | 36.750000 |
| 217292 | 2002-11-01 00:00:00 | 3.67 | 1.73 | 2.89 | NaN   | 0.38 | 113.900002 | 373.100006 | NaN  | 5.66  | 63.389999 |
| 217293 | 2002-11-01 00:00:00 | 1.37 | 0.58 | 1.17 | 2.37  | 0.15 | 65.389999  | 107.699997 | 1.30 | 9.11  | 9.640000  |
| 217294 | 2002-11-01 00:00:00 | 4.51 | 0.91 | 4.83 | 10.99 | NaN  | 149.800003 | 202.199997 | 1.00 | 5.75  | NaN       |
| 217295 | 2002-11-01 00:00:00 | 3.11 | 1.17 | 3.00 | 7.77  | 0.26 | 80.110001  | 180.300003 | 2.25 | 7.38  | 29.240000 |

217296 rows × 16 columns



```
In [81]: df=df.dropna()
```

```
In [82]: df=df.head(20)
```

```
In [83]: df.columns
```

```
Out[83]: Index(['date', 'BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',  
              'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station'],  
              dtype='object')
```

```
In [84]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 20 entries, 1 to 124  
Data columns (total 16 columns):  
#   Column      Non-Null Count  Dtype  
---  ---  
0   date        20 non-null    object  
1   BEN         20 non-null    float64  
2   CO          20 non-null    float64  
3   EBE         20 non-null    float64  
4   MXY         20 non-null    float64  
5   NMHC        20 non-null    float64  
6   NO_2        20 non-null    float64  
7   NOx         20 non-null    float64  
8   OXY         20 non-null    float64  
9   O_3         20 non-null    float64  
10  PM10        20 non-null    float64  
11  PXY         20 non-null    float64  
12  SO_2        20 non-null    float64  
13  TCH         20 non-null    float64  
14  TOL         20 non-null    float64  
15  station     20 non-null    int64  
dtypes: float64(14), int64(1), object(1)  
memory usage: 2.7+ KB
```


```
In [85]: data=df[['BEN', 'TOL', 'PXY']]
data
```

Out[85]:

|     | BEN  | TOL   | PXY  |
|-----|------|-------|------|
| 1   | 1.93 | 10.98 | 2.53 |
| 5   | 3.19 | 15.60 | 2.98 |
| 22  | 2.02 | 7.32  | 1.48 |
| 24  | 3.02 | 11.42 | 2.18 |
| 26  | 2.02 | 10.60 | 2.45 |
| 30  | 2.87 | 14.65 | 3.13 |
| 47  | 1.77 | 6.11  | 1.23 |
| 49  | 2.46 | 9.44  | 2.20 |
| 51  | 1.58 | 8.16  | 1.94 |
| 55  | 2.07 | 11.42 | 2.68 |
| 72  | 1.47 | 4.99  | 1.05 |
| 74  | 1.88 | 7.67  | 1.77 |
| 76  | 0.98 | 4.13  | 1.01 |
| 80  | 1.34 | 7.45  | 2.12 |
| 97  | 1.24 | 4.19  | 0.83 |
| 99  | 1.10 | 4.56  | 1.18 |
| 101 | 0.80 | 2.33  | 0.54 |
| 105 | 1.12 | 5.17  | 1.41 |
| 122 | 1.39 | 4.48  | 0.85 |
| 124 | 0.98 | 3.44  | 0.82 |


```
In [86]: data.plot.line(subplots=True)
```

Out[86]: array([<AxesSubplot:>, <AxesSubplot:>, <AxesSubplot:>], dtype=object)



```
In [87]: data.plot.line()
```


Out[87]: <AxesSubplot:>



```
In [88]: b=data[0:50]
```


```
In [89]: b.plot.bar()
```

Out[89]: <AxesSubplot:>




```
In [90]: data.plot.hist()
```

```
Out[90]: <AxesSubplot:ylabel='Frequency'>
```




```
In [91]: data.plot.area()
```

```
Out[91]: <AxesSubplot:>
```




```
In [92]: data.plot.box()
```

```
Out[92]: <AxesSubplot:>
```




```
In [93]: b.plot.pie(y='BEN' )
```

```
Out[93]: <AxesSubplot:ylabel='BEN'>
```



```
In [94]: data.plot.scatter(x='TOL', y='BEN')
```

```
Out[94]: <AxesSubplot:xlabel='TOL', ylabel='BEN'>
```



```
In [95]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 20 entries, 1 to 124
Data columns (total 16 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   date        20 non-null    object
 1   BEN         20 non-null    float64
 2   CO          20 non-null    float64
 3   EBE         20 non-null    float64
 4   MXY         20 non-null    float64
 5   NMHC        20 non-null    float64
 6   NO_2        20 non-null    float64
 7   NOx         20 non-null    float64
 8   OXY         20 non-null    float64
 9   O_3         20 non-null    float64
10   PM10        20 non-null    float64
11   PXY         20 non-null    float64
12   SO_2        20 non-null    float64
13   TCH         20 non-null    float64
14   TOL         20 non-null    float64
15   station     20 non-null    int64
dtypes: float64(14), int64(1), object(1)
memory usage: 2.7+ KB
```

In [96]: `df.describe()`

Out[96]:

|       | BEN       | CO        | EBE       | MXY       | NMHC      | NO_2       | NOx        | OXY       |
|-------|-----------|-----------|-----------|-----------|-----------|------------|------------|-----------|
| count | 20.000000 | 20.000000 | 20.000000 | 20.000000 | 20.000000 | 20.000000  | 20.000000  | 20.000000 |
| mean  | 1.761500  | 0.508000  | 1.781500  | 4.235500  | 0.094500  | 71.100000  | 96.450999  | 1.912500  |
| std   | 0.700115  | 0.273546  | 0.829358  | 2.019235  | 0.057809  | 26.411393  | 50.134432  | 0.892659  |
| min   | 0.800000  | 0.120000  | 0.490000  | 1.240000  | 0.000000  | 22.080000  | 24.139999  | 0.460000  |
| 25%   | 1.210000  | 0.272500  | 1.115000  | 2.477500  | 0.045000  | 44.017499  | 54.509998  | 1.257500  |
| 50%   | 1.675000  | 0.525000  | 1.660000  | 3.985000  | 0.115000  | 75.925003  | 88.530003  | 1.745000  |
| 75%   | 2.032500  | 0.712500  | 2.342500  | 5.557500  | 0.130000  | 92.164999  | 136.874996 | 2.580000  |
| max   | 3.190000  | 1.040000  | 3.450000  | 7.920000  | 0.210000  | 113.699997 | 195.399994 | 3.730000  |

In [97]: `df1=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3', 'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station']]`

In [\*]: `sns.pairplot(df1[0:50])`

Out[98]: `<seaborn.axisgrid.PairGrid at 0x19113de65b0>`

In [\*]: `sns.distplot(df1['station'])`

In [\*]: `sns.heatmap(df1.corr())`

In [\*]: `x=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3', 'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]`  
`y=df['station']`

In [\*]: `from sklearn.model_selection import train_test_split`  
`x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)`

In [\*]: `from sklearn.linear_model import LinearRegression`  
`lr=LinearRegression()`  
`lr.fit(x_train,y_train)`

In [\*]: `lr.intercept_`

In [\*]: `coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])`  
`coeff`



```
In [*]: prediction =lr.predict(x_test)
plt.scatter(y_test,prediction)
```

```
In [*]: lr.score(x_test,y_test)
```

```
In [*]: lr.score(x_train,y_train)
```

```
In [*]: from sklearn.linear_model import Ridge,Lasso
```

```
In [*]: rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
```

```
In [*]: rr.score(x_test,y_test)
```

```
In [*]: rr.score(x_train,y_train)
```

```
In [*]: la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

```
In [*]: la.score(x_train,y_train)
```

```
In [*]: la.score(x_test,y_test)
```

```
In [*]: from sklearn.linear_model import ElasticNet
en=ElasticNet()
en.fit(x_train,y_train)
```

```
In [*]: en.coef_
```

```
In [*]: en.intercept_
```

```
In [*]: prediction=en.predict(x_test)
```

```
In [*]: en.score(x_test,y_test)
```

```
In [*]: from sklearn import metrics
print(metrics.mean_absolute_error(y_test,prediction))
print(metrics.mean_squared_error(y_test,prediction))
print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

```
In [*]: from sklearn.linear_model import LogisticRegression
```

```
In [*]: feature_matrix=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_
      'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]
target_vector=df[ 'station']
```

```
In [*]: feature_matrix.shape
```

```
In [56]: target_vector.shape
```

```
Out[56]: (20,)
```

```
In [57]: from sklearn.preprocessing import StandardScaler
```

```
In [58]: fs=StandardScaler().fit_transform(feature_matrix)
```

```
In [59]: logr=LogisticRegression(max_iter=10000)
logr.fit(fs,target_vector)
```

```
Out[59]: LogisticRegression(max_iter=10000)
```

```
In [60]: observation=[[1,2,3,4,5,6,7,8,9,10,11,12,13,14]]
```

```
In [61]: prediction=logr.predict(observation)
print(prediction)

[28079099]
```

```
In [62]: logr.score(fs,target_vector)
```

```
Out[62]: 1.0
```

```
In [63]: logr.predict_proba(observation)[0][0]
```

```
Out[63]: 0.021267025366369423
```

```
In [64]: logr.predict_proba(observation)
```

```
Out[64]: array([[2.12670254e-02, 7.21980940e-16, 4.47700636e-08, 9.78732930e-01]])
```

```
In [65]: from sklearn.ensemble import RandomForestClassifier
```

```
In [66]: rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
```

```
Out[66]: RandomForestClassifier()
```

```
In [67]: parameters={'max_depth':[1,2,3,4,5],
'min_samples_leaf':[5,10,15,20,25],
'n_estimators':[10,20,30,40,50]}
```

```
In [68]: from sklearn.model_selection import GridSearchCV
grid_search =GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="acc
grid_search.fit(x_train,y_train)
```

```
Out[68]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
param_grid={'max_depth': [1, 2, 3, 4, 5],
'min_samples_leaf': [5, 10, 15, 20, 25],
'n_estimators': [10, 20, 30, 40, 50]},
scoring='accuracy')
```

```
In [69]: grid_search.best_score_
```

```
Out[69]: 0.3571428571428571
```

```
In [70]: rfc_best=grid_search.best_estimator_
```

```
In [74]: from sklearn.tree import plot_tree
plt.figure(figsize=(50,5))
plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names=['a','b'],
```

```
Out[74]: [Text(1395.0, 135.9, 'gini = 0.612\nsamples = 8\nvalue = [1, 1, 7, 5]\nclass
= c')]
```

```
gini = 0.612
samples = 8
value = [1, 1, 7, 5]
class = c
```

## Conclusion

Linear Regression =1.0

Ridge Regression =0.6834711491207041

Lasso Regression =0.6132115096399764

ElasticNet Regression =-2.6934722951520897

Logistic Regression =-2.6934722951520897

Randomforest =-2.69347229515208971

Logistic Regression is suitable for this dataset

In [ ]: