

03/08/2023 (P6)

```
In [161]: import numpy as np  
import pandas as pd  
import seaborn as sns  
import matplotlib.pyplot as plt
```

```
In [163]: df=pd.read_csv(r"C:\Users\user\Downloads\madrid_2006.csv")
df
```

Out[163]:

	date	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	O_3	PM10	PM25	PXY	SO_2	TCH
0	2006-02-01 01:00:00	NaN	1.84	NaN	NaN	NaN	155.100006	490.100006	NaN	4.880000	97.570000	40.259998	NaN	33.779999	NaN
1	2006-02-01 01:00:00	1.68	1.01	2.38	6.36	0.32	94.339996	229.699997	3.04	7.100000	25.820000	NaN	2.48	11.890000	1.59
2	2006-02-01 01:00:00	NaN	1.25	NaN	NaN	NaN	66.800003	192.000000	NaN	4.430000	34.419998	NaN	NaN	19.719999	NaN
3	2006-02-01 01:00:00	NaN	1.68	NaN	NaN	NaN	103.000000	407.799988	NaN	4.830000	28.260000	NaN	NaN	21.129999	NaN
4	2006-02-01 01:00:00	NaN	1.31	NaN	NaN	NaN	105.400002	269.200012	NaN	6.990000	54.180000	NaN	NaN	11.050000	NaN
...
230563	2006-05-01 00:00:00	5.88	0.83	6.23	NaN	0.20	112.500000	218.000000	NaN	24.389999	93.120003	NaN	NaN	7.400000	1.50
230564	2006-05-01 00:00:00	0.76	0.32	0.48	1.09	0.08	51.900002	54.820000	0.61	48.410000	29.469999	15.640000	0.50	8.840000	1.32
230565	2006-05-01 00:00:00	0.96	NaN	0.69	NaN	0.19	135.100006	179.199997	NaN	11.460000	64.680000	35.000000	NaN	12.110000	1.51
230566	2006-05-01 00:00:00	0.50	NaN	0.67	NaN	0.10	82.599998	105.599998	NaN	NaN	94.360001	NaN	NaN	4.890000	1.47
230567	2006-05-01 00:00:00	1.95	0.74	1.99	4.00	0.24	107.300003	160.199997	2.01	17.730000	52.490002	27.920000	1.70	9.170000	1.50

230568 rows × 17 columns



```
In [169]: df=df.dropna()
```

```
In [170]: df.columns
```

```
Out[170]: Index(['date', 'BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
                  'PM10', 'PM25', 'PXY', 'SO_2', 'TCH', 'TOL', 'station'],
                  dtype='object')
```

```
In [171]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 24758 entries, 5 to 230567
Data columns (total 17 columns):
 #   Column    Non-Null Count  Dtype  
--- 
 0   date      24758 non-null   object 
 1   BEN       24758 non-null   float64
 2   CO        24758 non-null   float64
 3   EBE       24758 non-null   float64
 4   MXY       24758 non-null   float64
 5   NMHC      24758 non-null   float64
 6   NO_2      24758 non-null   float64
 7   NOx       24758 non-null   float64
 8   OXY       24758 non-null   float64
 9   O_3        24758 non-null   float64
 10  PM10      24758 non-null   float64
 11  PM25      24758 non-null   float64
 12  PXY       24758 non-null   float64
 13  SO_2      24758 non-null   float64
 14  TCH       24758 non-null   float64
 15  TOL       24758 non-null   float64
 16  station    24758 non-null   int64  
dtypes: float64(15), int64(1), object(1)
memory usage: 3.4+ MB
```

```
In [172]: data=df[['EBE', 'MXY', 'PXY']]  
data
```

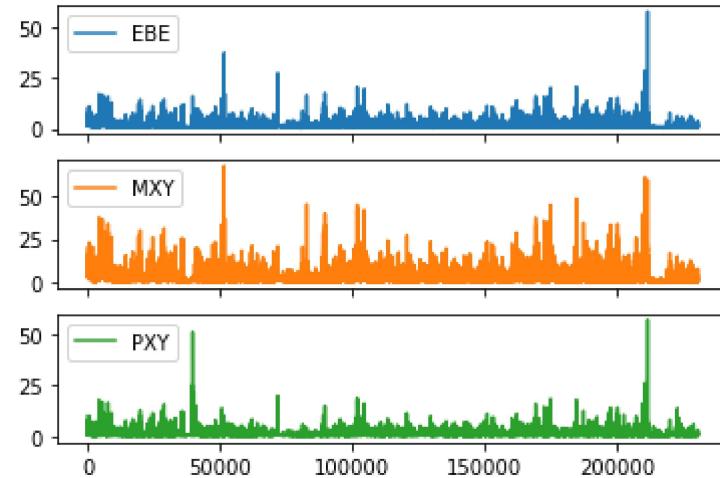
Out[172]:

	EBE	MXY	PXY
5	9.98	19.959999	10.11
22	1.24	2.670000	0.79
25	2.64	9.660000	4.46
31	7.92	17.139999	8.06
48	1.24	2.740000	0.82
...
230538	0.37	0.430000	1.00
230541	1.53	2.200000	1.21
230547	3.71	7.960000	3.36
230564	0.48	1.090000	0.50
230567	1.99	4.000000	1.70

24758 rows × 3 columns

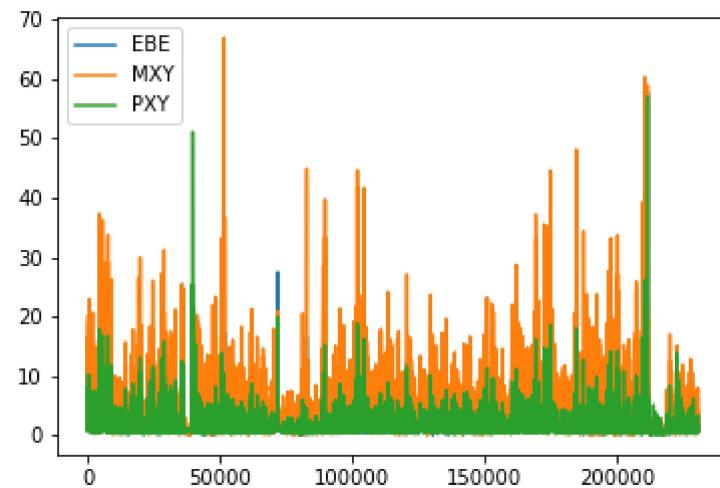
```
In [173]: data.plot.line(subplots=True)
```

```
Out[173]: array([<AxesSubplot:>, <AxesSubplot:>, <AxesSubplot:>], dtype=object)
```



```
In [174]: data.plot.line()
```

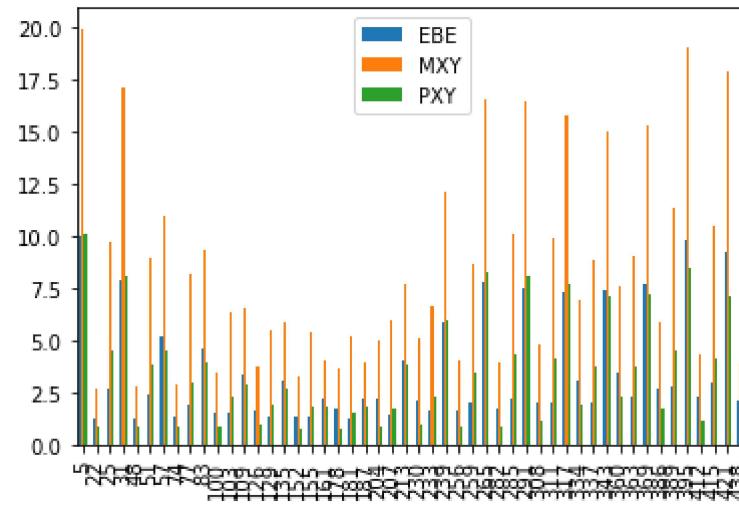
```
Out[174]: <AxesSubplot:>
```



```
In [175]: b=data[0:50]
```

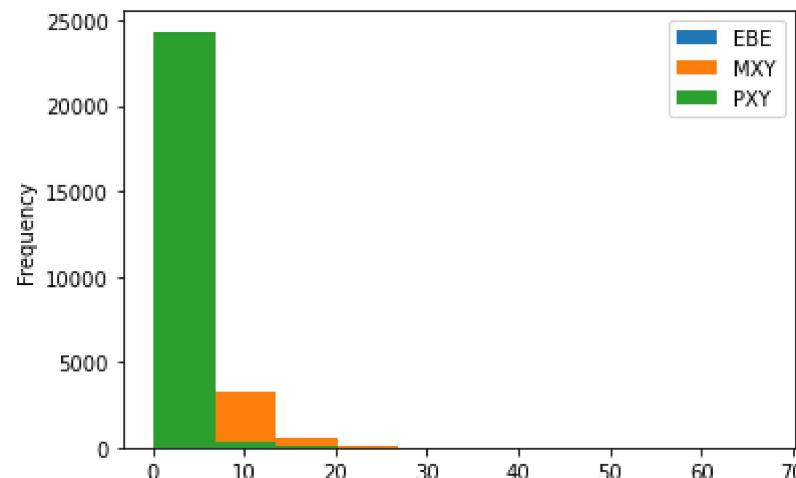
```
In [176]: b.plot.bar()
```

Out[176]: <AxesSubplot:>



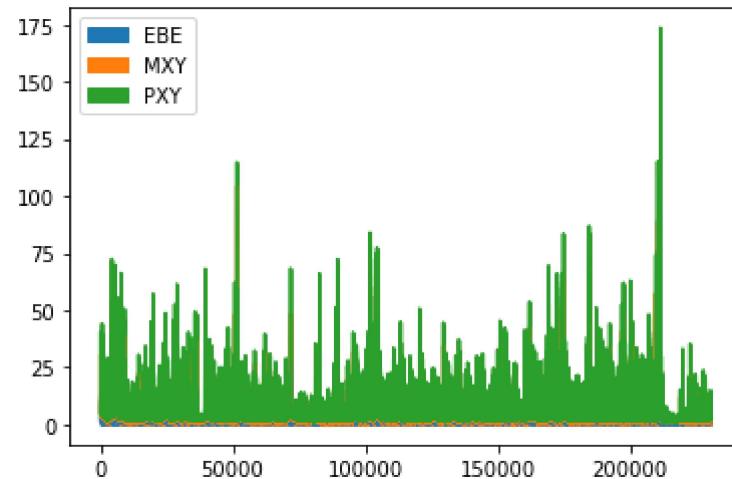
```
In [177]: data.plot.hist()
```

Out[177]: <AxesSubplot:ylabel='Frequency'>



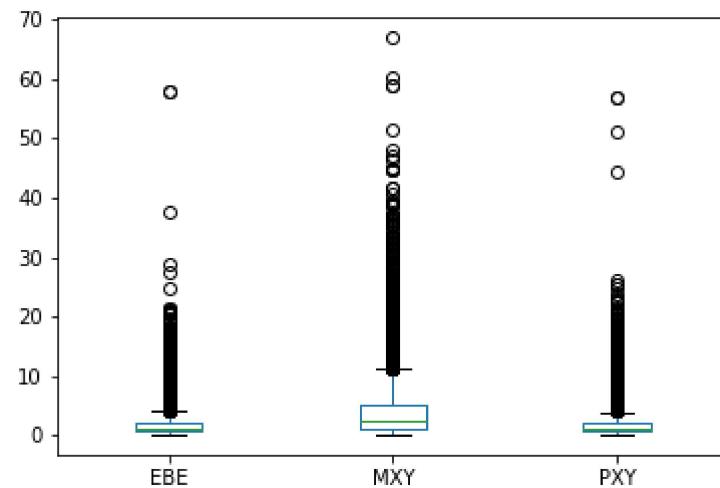
```
In [178]: data.plot.area()
```

```
Out[178]: <AxesSubplot:>
```



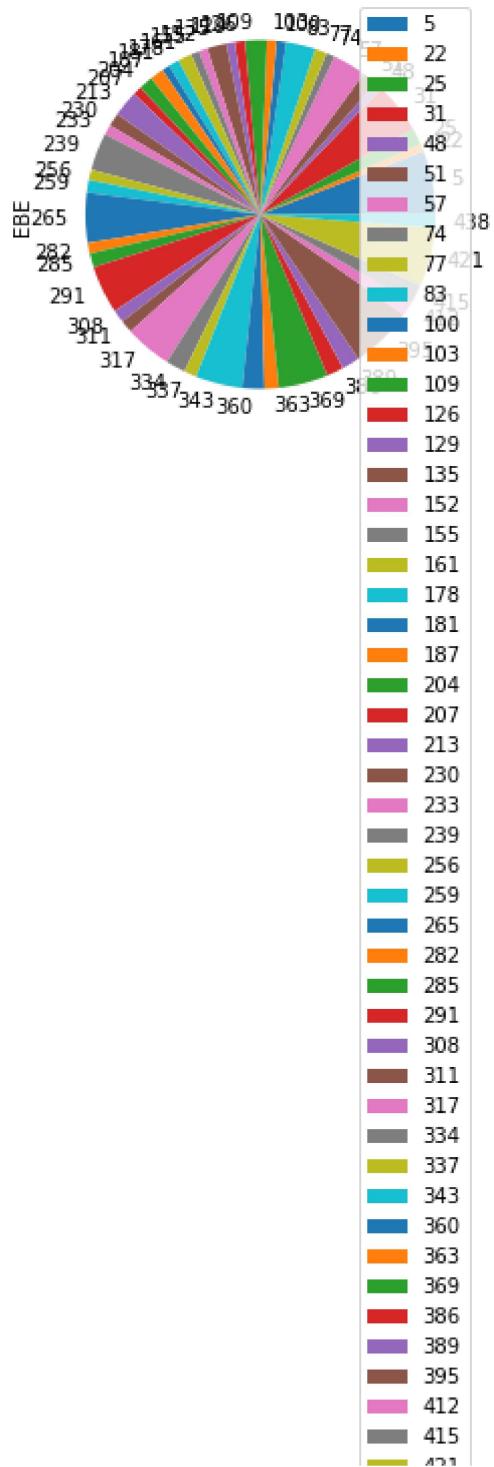
```
In [179]: data.plot.box()
```

```
Out[179]: <AxesSubplot:>
```



```
In [180]: b.plot.pie(y='EBE' )
```

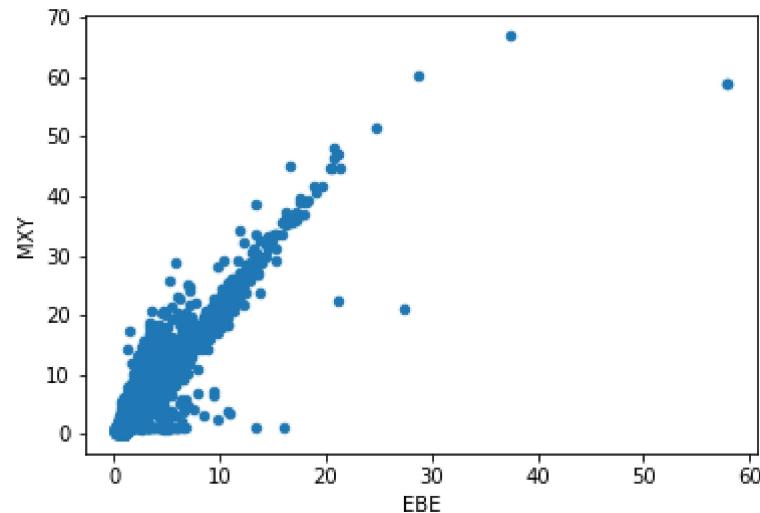
```
Out[180]: <AxesSubplot:ylabel='EBE'>
```



```
In [181]: data.plot.scatter(x='EBE' ,y='MXY')
```

```
Out[181]: <AxesSubplot:xlabel='EBE', ylabel='MXY'>
```



In [182]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 24758 entries, 5 to 230567
Data columns (total 17 columns):
 #   Column    Non-Null Count  Dtype  
--- 
 0   date      24758 non-null   object 
 1   BEN       24758 non-null   float64
 2   CO        24758 non-null   float64
 3   EBE       24758 non-null   float64
 4   MXY       24758 non-null   float64
 5   NMHC      24758 non-null   float64
 6   NO_2      24758 non-null   float64
 7   NOx       24758 non-null   float64
 8   OXY       24758 non-null   float64
 9   O_3        24758 non-null   float64
 10  PM10      24758 non-null   float64
 11  PM25      24758 non-null   float64
 12  PXY       24758 non-null   float64
 13  SO_2      24758 non-null   float64
 14  TCH       24758 non-null   float64
 15  TOL       24758 non-null   float64
 16  station   24758 non-null   int64  
dtypes: float64(15), int64(1), object(1)
memory usage: 3.4+ MB
```

```
In [183]: df.describe()
```

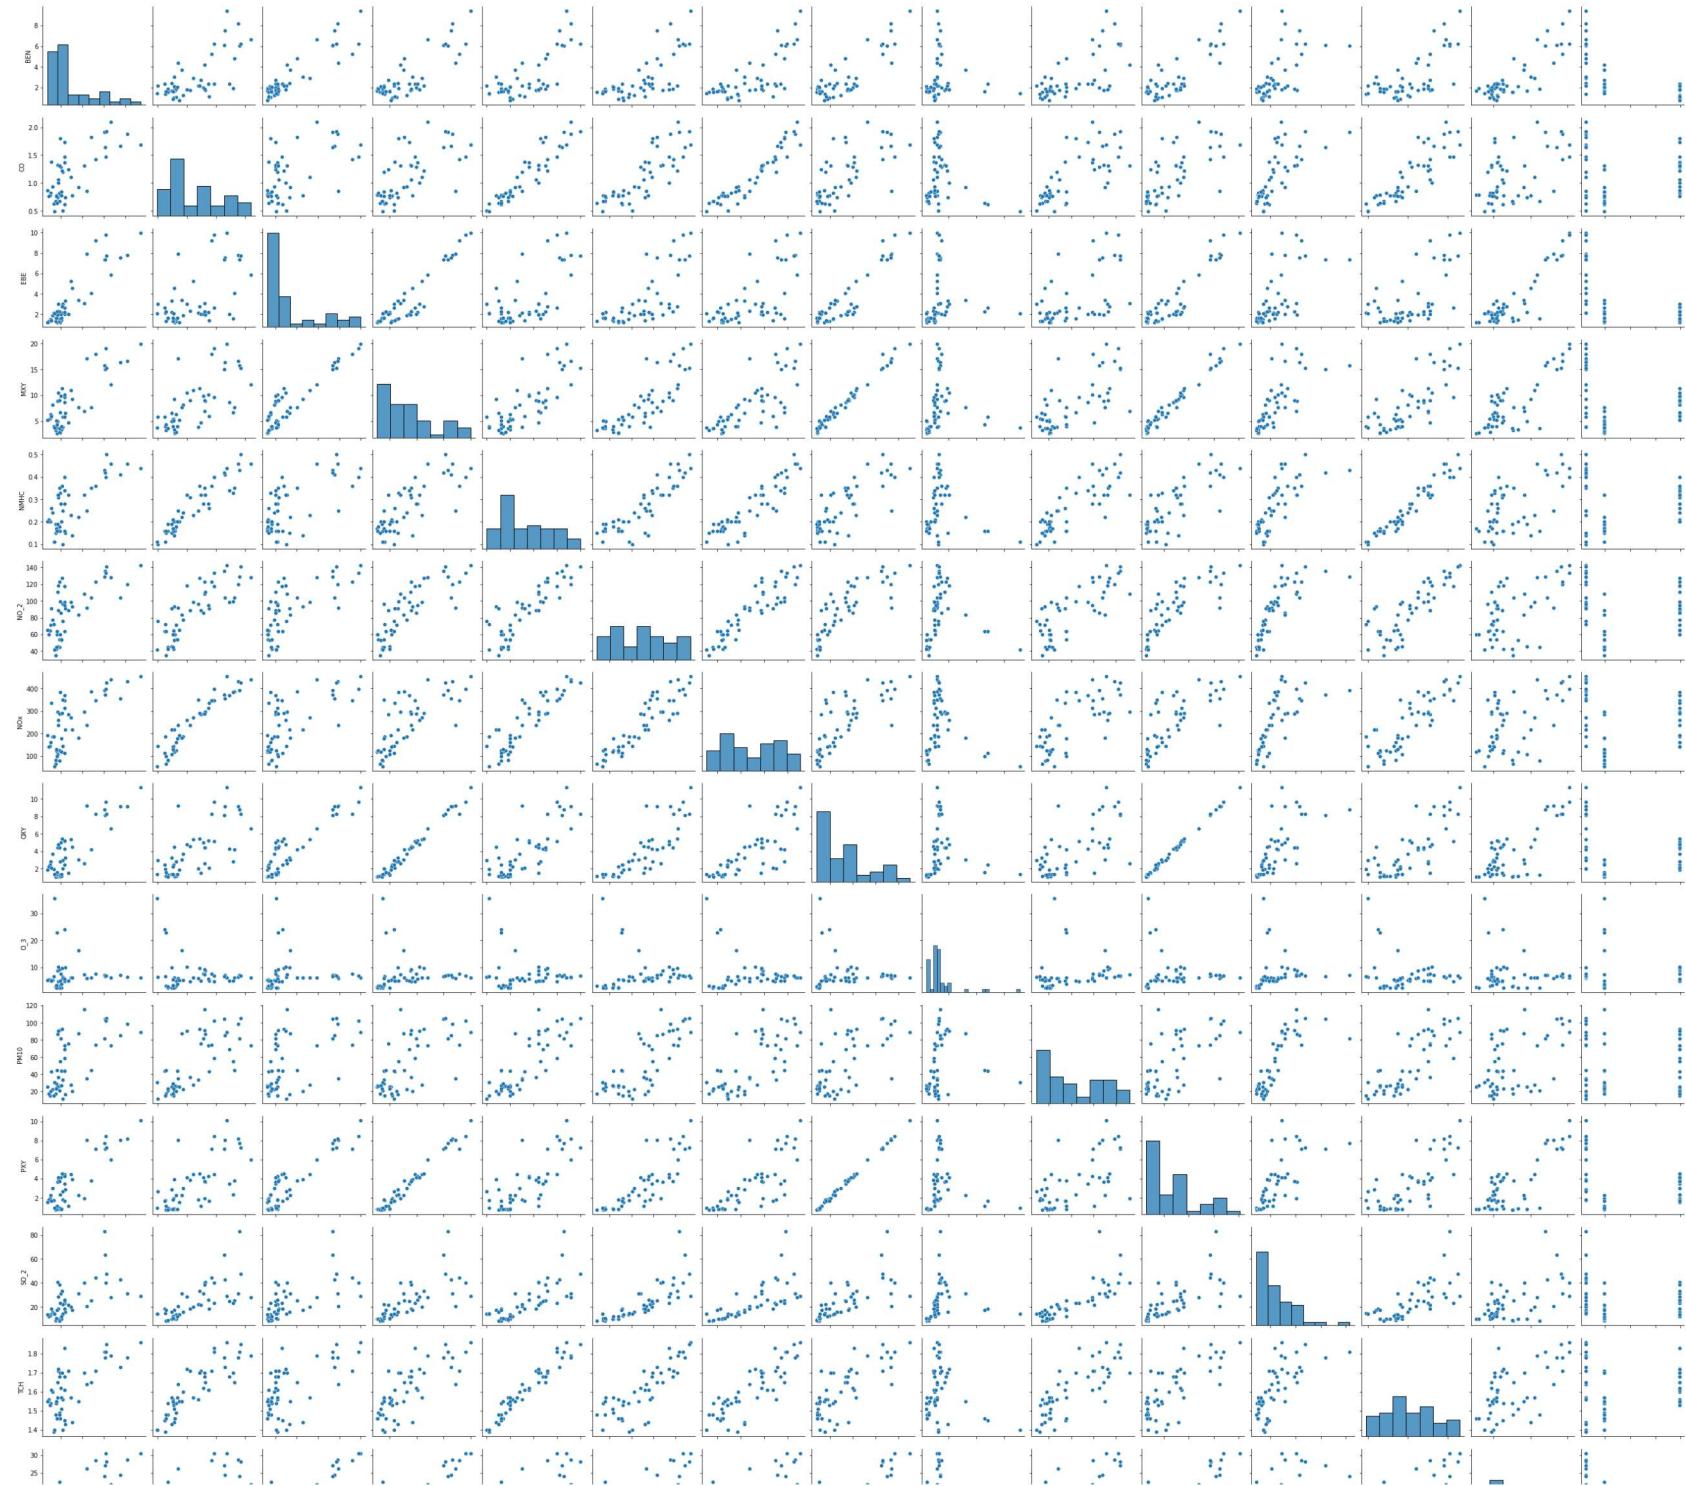
Out[183]:

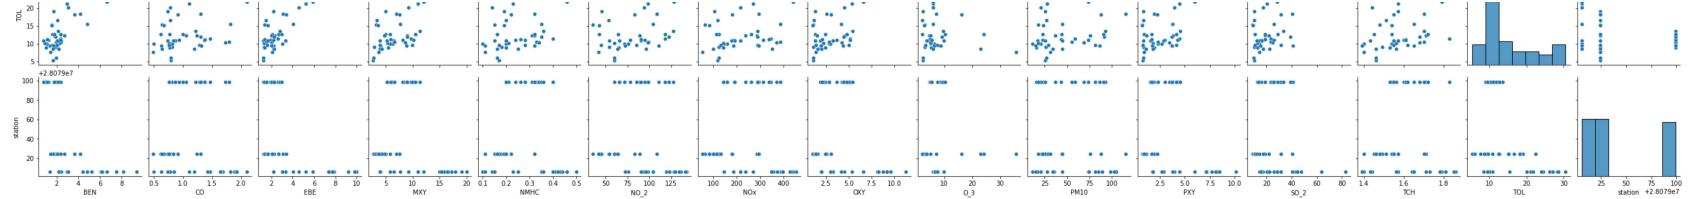
	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	O_3
count	24758.000000	24758.000000	24758.000000	24758.000000	24758.000000	24758.000000	24758.000000	24758.000000	24758.000000
mean	1.350624	0.600713	1.824534	3.835034	0.176546	58.333481	116.419090	1.990347	37.155685
std	1.541636	0.419048	1.868939	4.069036	0.126683	40.529382	117.557064	1.931620	31.127681
min	0.110000	0.000000	0.170000	0.150000	0.000000	1.680000	2.020000	0.190000	0.310000
25%	0.450000	0.360000	0.810000	1.060000	0.100000	28.450001	36.882501	0.960000	12.110000
50%	0.850000	0.500000	1.130000	2.500000	0.150000	52.959999	85.180000	1.260000	28.675000
75%	1.680000	0.720000	2.160000	5.090000	0.220000	79.347498	158.300003	2.470000	53.029999
max	45.430000	7.250000	57.799999	66.900002	2.020000	461.299988	1680.000000	63.000000	178.899994

```
In [184]: df1=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',  
'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station']]
```

```
In [185]: sns.pairplot(df1[0:50])
```

```
Out[185]: <seaborn.axisgrid.PairGrid at 0x1b170cdca30>
```

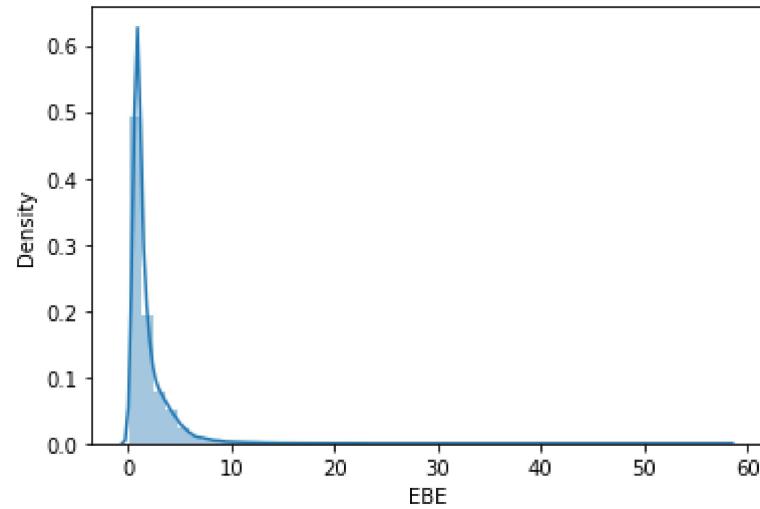





```
In [186]: sns.distplot(df1['EBE'])
```

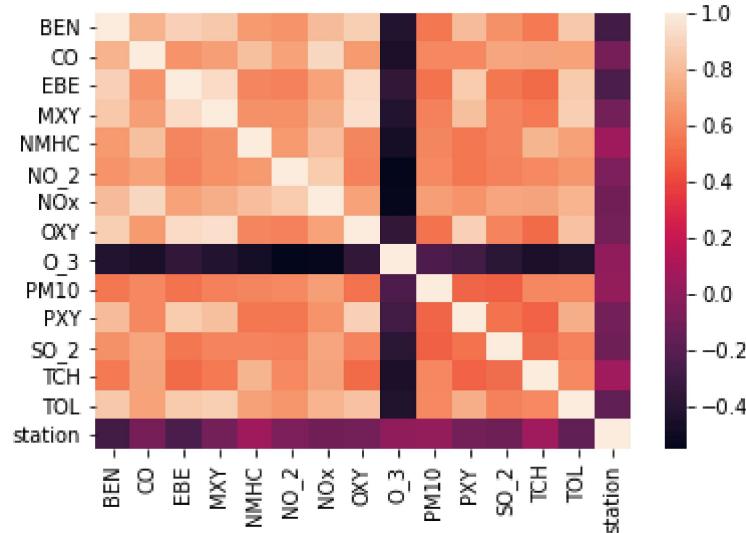
C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
warnings.warn(msg, FutureWarning)

```
Out[186]: <AxesSubplot:xlabel='EBE', ylabel='Density'>
```



```
In [187]: sns.heatmap(df1.corr())
```

```
Out[187]: <AxesSubplot:>
```



```
In [188]: x=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',  
          'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]  
y=df['station']
```

```
In [189]: from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [190]: from sklearn.linear_model import LinearRegression  
lr=LinearRegression()  
lr.fit(x_train,y_train)
```

```
Out[190]: LinearRegression()
```

```
In [191]: lr.intercept_
```

```
Out[191]: 28079026.2561639
```

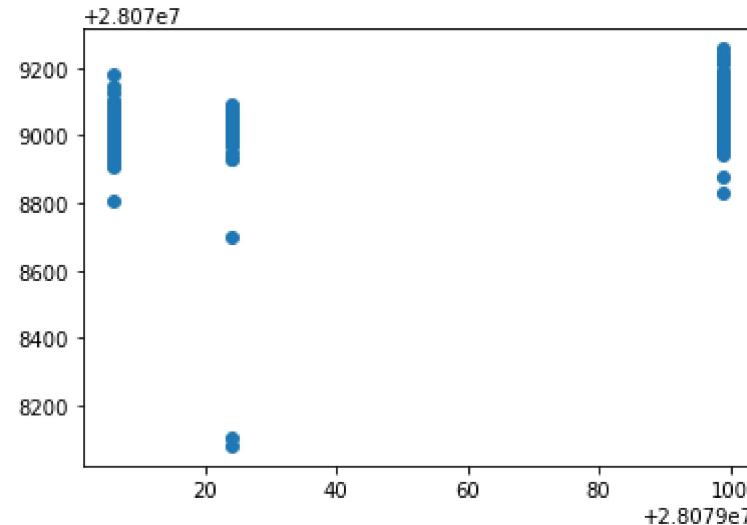
```
In [192]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=[ 'Co-efficient'])
coeff
```

Out[192]:

	Co-efficient
BEN	-21.893539
CO	-7.439464
EBE	-24.811502
MXY	6.532366
NMHC	127.075045
NO_2	-0.042926
NOx	0.019741
OXY	10.767984
O_3	-0.037720
PM10	0.118907
PXY	6.860635
SO_2	-0.512966
TCH	13.701935
TOL	-0.018435

```
In [193]: prediction = lr.predict(x_test)
plt.scatter(y_test,prediction)
```

```
Out[193]: <matplotlib.collections.PathCollection at 0x1b17ee8c3d0>
```



```
In [194]: lr.score(x_test,y_test)
```

```
Out[194]: 0.2762628434964486
```

```
In [195]: lr.score(x_train,y_train)
```

```
Out[195]: 0.4181657194466114
```

```
In [196]: from sklearn.linear_model import Ridge,Lasso
```

```
In [197]: rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
```

```
Out[197]: Ridge(alpha=10)
```

```
In [198]: rr.score(x_test,y_test)
```

```
Out[198]: 0.27378982126458806
```

```
In [199]: rr.score(x_train,y_train)
```

```
Out[199]: 0.41752121857497326
```

```
In [200]: la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

```
Out[200]: Lasso(alpha=10)
```

```
In [201]: la.score(x_train,y_train)
```

```
Out[201]: 0.062109979229048484
```

```
In [202]: la.score(x_test,y_test)
```

```
Out[202]: 0.05849392873330106
```

```
In [203]: from sklearn.linear_model import ElasticNet
en=ElasticNet()
en.fit(x_train,y_train)
```

```
Out[203]: ElasticNet()
```

```
In [204]: en.coef_
```

```
Out[204]: array([-9.39800478,  0.          , -9.38981297,  3.68157567,  0.42838989,
   -0.01386303,  0.01633955,  2.57991111, -0.11699899,  0.29745146,
   2.3894965 , -0.39224048,  0.52916397, -0.854195  ])
```

```
In [205]: en.intercept_
```

```
Out[205]: 28079052.203664407
```

```
In [206]: prediction=en.predict(x_test)
```

```
In [207]: en.score(x_test,y_test)
```

```
Out[207]: 0.220680994144546
```

```
In [208]: from sklearn import metrics  
print(metrics.mean_absolute_error(y_test,prediction))  
print(metrics.mean_squared_error(y_test,prediction))  
print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

```
32.08977616142971  
1285.8284012148508  
35.85844950935345
```

```
In [209]: from sklearn.linear_model import LogisticRegression
```

```
In [210]: feature_matrix=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',  
'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]  
target_vector=df[ 'station']
```

```
In [211]: feature_matrix.shape
```

```
Out[211]: (24758, 14)
```

```
In [212]: target_vector.shape
```

```
Out[212]: (24758,)
```

```
In [213]: from sklearn.preprocessing import StandardScaler
```

```
In [214]: fs=StandardScaler().fit_transform(feature_matrix)
```

```
In [215]: logr=LogisticRegression(max_iter=10000)  
logr.fit(fs,target_vector)
```

```
Out[215]: LogisticRegression(max_iter=10000)
```

```
In [216]: observation=[[1,2,3,4,5,6,7,8,9,10,11,12,13,14]]
```

```
In [217]: prediction=logr.predict(observation)  
print(prediction)
```

```
[28079099]
```

```
In [218]: logr.classes_
```

```
Out[218]: array([28079006, 28079024, 28079099], dtype=int64)
```

```
In [219]: logr.score(fs,target_vector)
```

```
Out[219]: 0.8741416915744405
```

```
In [220]: logr.predict_proba(observation)[0][0]
```

```
Out[220]: 3.5557727473608076e-15
```

```
In [221]: logr.predict_proba(observation)
```

```
Out[221]: array([[3.55577275e-15, 7.80743173e-29, 1.00000000e+00]])
```

```
In [222]: from sklearn.ensemble import RandomForestClassifier
```

```
In [223]: rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
```

```
Out[223]: RandomForestClassifier()
```

```
In [224]: parameters={'max_depth':[1,2,3,4,5],
 'min_samples_leaf':[5,10,15,20,25],
 'n_estimators':[10,20,30,40,50]
 }
```

```
In [225]: from sklearn.model_selection import GridSearchCV
grid_search =GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="accuracy")
grid_search.fit(x_train,y_train)
```

```
Out[225]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
 param_grid={'max_depth': [1, 2, 3, 4, 5],
 'min_samples_leaf': [5, 10, 15, 20, 25],
 'n_estimators': [10, 20, 30, 40, 50]},
 scoring='accuracy')
```

```
In [226]: grid_search.best_score_
```

```
Out[226]: 0.8745527986151183
```

```
In [227]: rfc_best=grid_search.best_estimator_
```

```
In [228]: from sklearn.tree import plot_tree
plt.figure(figsize=(80,40))
plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names=['a','b','c','d'],filled=True)
```

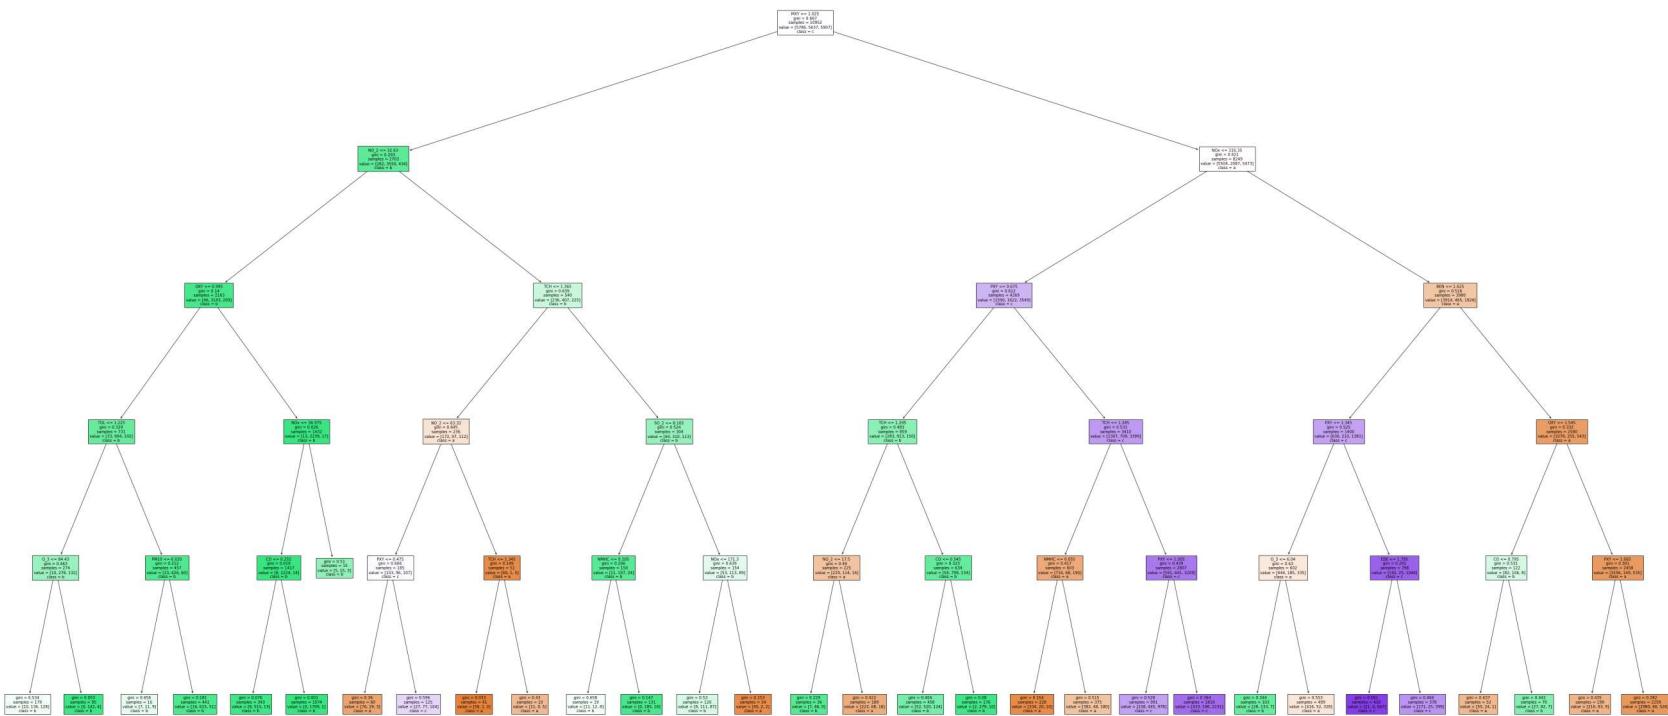
Out[228]: [Text(2148.3, 1993.2, 'MXY <= 1.025\ngini = 0.667\nsamples = 10952\nvalue = [5786, 5637, 5907]\nnclass = c'),
Text(1023.000000000001, 1630.800000000002, 'NO_2 <= 32.63\ngini = 0.293\nsamples = 2703\nvalue = [282, 35
50, 434]\nnclass = b'),
Text(558.0, 1268.4, 'OXY <= 0.995\ngini = 0.14\nsamples = 2163\nvalue = [46, 3143, 209]\nnclass = b'),
Text(297.6, 906.0, 'TOL <= 1.225\ngini = 0.329\nsamples = 731\nvalue = [33, 904, 192]\nnclass = b'),
Text(148.8, 543.599999999999, 'O_3 <= 84.43\ngini = 0.463\nsamples = 274\nvalue = [10, 278, 132]\nnclass =
b'),
Text(74.4, 181.1999999999982, 'gini = 0.534\nsamples = 179\nvalue = [10, 136, 128]\nnclass = b'),
Text(223.200000000002, 181.1999999999982, 'gini = 0.053\nsamples = 95\nvalue = [0, 142, 4]\nnclass = b'),
Text(446.400000000003, 543.599999999999, 'PM10 <= 6.035\ngini = 0.212\nsamples = 457\nvalue = [23, 626,
60]\nnclass = b'),
Text(372.0, 181.1999999999982, 'gini = 0.656\nsamples = 16\nvalue = [7, 11, 9]\nnclass = b'),
Text(520.800000000001, 181.1999999999982, 'gini = 0.181\nsamples = 441\nvalue = [16, 615, 51]\nnclass =
b'),
Text(818.400000000001, 906.0, 'NOx <= 36.975\ngini = 0.026\nsamples = 1432\nvalue = [13, 2239, 17]\nnclass
= b'),
Text(744.0, 543.599999999999, 'CO <= 0.255\ngini = 0.019\nsamples = 1417\nvalue = [8, 2224, 14]\nnclass =
b'),
Text(669.6, 181.1999999999982, 'gini = 0.076\nsamples = 343\nvalue = [8, 515, 13]\nnclass = b'),
Text(818.400000000001, 181.1999999999982, 'gini = 0.001\nsamples = 1074\nvalue = [0, 1709, 1]\nnclass =
b'),
Text(892.800000000001, 543.599999999999, 'gini = 0.51\nsamples = 15\nvalue = [5, 15, 3]\nnclass = b'),
Text(1488.0, 1268.4, 'TCH <= 1.365\ngini = 0.639\nsamples = 540\nvalue = [236, 407, 225]\nnclass = b'),
Text(1190.4, 906.0, 'NO_2 <= 63.32\ngini = 0.645\nsamples = 236\nvalue = [172, 97, 112]\nnclass = a'),
Text(1041.600000000001, 543.599999999999, 'PXY <= 0.475\ngini = 0.666\nsamples = 185\nvalue = [103, 96, 1
07]\nnclass = c'),
Text(967.2, 181.1999999999982, 'gini = 0.36\nsamples = 60\nvalue = [76, 19, 3]\nnclass = a'),
Text(1116.0, 181.1999999999982, 'gini = 0.596\nsamples = 125\nvalue = [27, 77, 104]\nnclass = c'),
Text(1339.2, 543.599999999999, 'TCH <= 1.345\ngini = 0.149\nsamples = 51\nvalue = [69, 1, 5]\nnclass = a'),
Text(1264.800000000002, 181.1999999999982, 'gini = 0.033\nsamples = 41\nvalue = [58, 1, 0]\nnclass = a'),
Text(1413.600000000001, 181.1999999999982, 'gini = 0.43\nsamples = 10\nvalue = [11, 0, 5]\nnclass = a'),
Text(1785.600000000001, 906.0, 'SO_2 <= 8.165\ngini = 0.524\nsamples = 304\nvalue = [64, 310, 113]\nnclass
= b'),
Text(1636.800000000002, 543.599999999999, 'NMHC <= 0.105\ngini = 0.266\nsamples = 150\nvalue = [11, 197,
24]\nnclass = b'),
Text(1562.4, 181.1999999999982, 'gini = 0.658\nsamples = 19\nvalue = [11, 12, 8]\nnclass = b'),
Text(1711.2, 181.1999999999982, 'gini = 0.147\nsamples = 131\nvalue = [0, 185, 16]\nnclass = b'),
Text(1934.4, 543.599999999999, 'NOx <= 171.3\ngini = 0.639\nsamples = 154\nvalue = [53, 113, 89]\nnclass =
b'),
Text(1860.000000000002, 181.1999999999982, 'gini = 0.53\nsamples = 120\nvalue = [8, 111, 87]\nnclass =
b'),
Text(2008.800000000002, 181.1999999999982, 'gini = 0.153\nsamples = 34\nvalue = [45, 2, 2]\nnclass = a'),
Text(3273.600000000004, 1630.800000000002, 'NOx <= 116.35\ngini = 0.621\nsamples = 8249\nvalue = [5504, 2

087, 5473]\nclass = a'),
Text(2678.4, 1268.4, 'PXY <= 0.675\ngini = 0.612\nsamples = 4269\nvalue = [1590, 1622, 3549]\nclass = c'),
Text(2380.8, 906.0, 'TCH <= 1.295\ngini = 0.483\nsamples = 859\nvalue = [283, 913, 150]\nclass = b'),
Text(2232.0, 543.5999999999999, 'NO_2 <= 17.5\ngini = 0.49\nsamples = 225\nvalue = [229, 114, 16]\nclass = a'),
Text(2157.600000000004, 181.1999999999982, 'gini = 0.229\nsamples = 36\nvalue = [7, 46, 0]\nclass = b'),
Text(2306.4, 181.1999999999982, 'gini = 0.422\nsamples = 189\nvalue = [222, 68, 16]\nclass = a'),
Text(2529.600000000004, 543.5999999999999, 'CO <= 0.545\ngini = 0.323\nsamples = 634\nvalue = [54, 799, 134]\nclass = b'),
Text(2455.200000000003, 181.1999999999982, 'gini = 0.404\nsamples = 458\nvalue = [52, 520, 124]\nclass = b'),
Text(2604.0, 181.1999999999982, 'gini = 0.08\nsamples = 176\nvalue = [2, 279, 10]\nclass = b'),
Text(2976.0, 906.0, 'TCH <= 1.285\ngini = 0.531\nsamples = 3410\nvalue = [1307, 709, 3399]\nclass = c'),
Text(2827.200000000003, 543.5999999999999, 'NMHC <= 0.055\ngini = 0.417\nsamples = 603\nvalue = [716, 68, 190]\nclass = a'),
Text(2752.8, 181.1999999999982, 'gini = 0.154\nsamples = 228\nvalue = [334, 20, 10]\nclass = a'),
Text(2901.600000000004, 181.1999999999982, 'gini = 0.515\nsamples = 375\nvalue = [382, 48, 180]\nclass = a'),
Text(3124.8, 543.5999999999999, 'PXY <= 1.005\ngini = 0.439\nsamples = 2807\nvalue = [591, 641, 3209]\nclass = c'),
Text(3050.4, 181.1999999999982, 'gini = 0.528\nsamples = 991\nvalue = [158, 445, 978]\nclass = c'),
Text(3199.200000000003, 181.1999999999982, 'gini = 0.364\nsamples = 1816\nvalue = [433, 196, 2231]\nclass = c'),
Text(3868.8, 1268.4, 'BEN <= 1.625\ngini = 0.516\nsamples = 3980\nvalue = [3914, 465, 1924]\nclass = a'),
Text(3571.200000000003, 906.0, 'PXY <= 1.345\ngini = 0.525\nsamples = 1400\nvalue = [636, 210, 1381]\nclass = c'),
Text(3422.4, 543.5999999999999, 'O_3 <= 6.04\ngini = 0.63\nsamples = 602\nvalue = [444, 185, 335]\nclass = a'),
Text(3348.000000000005, 181.1999999999982, 'gini = 0.344\nsamples = 103\nvalue = [28, 133, 7]\nclass = b'),
Text(3496.8, 181.1999999999982, 'gini = 0.553\nsamples = 499\nvalue = [416, 52, 328]\nclass = a'),
Text(3720.000000000005, 543.5999999999999, 'EBE <= 1.795\ngini = 0.291\nsamples = 798\nvalue = [192, 25, 1046]\nclass = c'),
Text(3645.600000000004, 181.1999999999982, 'gini = 0.061\nsamples = 420\nvalue = [21, 0, 647]\nclass = c'),
Text(3794.4, 181.1999999999982, 'gini = 0.466\nsamples = 378\nvalue = [171, 25, 399]\nclass = c'),
Text(4166.400000000001, 906.0, 'OXY <= 1.545\ngini = 0.332\nsamples = 2580\nvalue = [3278, 255, 543]\nclass = a'),
Text(4017.600000000004, 543.5999999999999, 'CO <= 0.795\ngini = 0.531\nsamples = 122\nvalue = [82, 106, 8]\nclass = b'),
Text(3943.200000000003, 181.1999999999982, 'gini = 0.437\nsamples = 52\nvalue = [55, 24, 1]\nclass = a'),
Text(4092.000000000005, 181.1999999999982, 'gini = 0.442\nsamples = 70\nvalue = [27, 82, 7]\nclass = b'),
Text(4315.200000000001, 543.5999999999999, 'PXY <= 1.665\ngini = 0.301\nsamples = 2458\nvalue = [3196, 149,

```

535]\nclass = a'),
Text(4240.8, 181.1999999999982, 'gini = 0.435\ncounts = 199\nvalue = [216, 83, 9]\nclass = a'),
Text(4389.6, 181.1999999999982, 'gini = 0.282\ncounts = 2259\nvalue = [2980, 66, 526]\nclass = a')]

```



Conclusion

```
In [ ]: linear Regression=0.4181657194466114  
Ridge Regression=0.41752121857497326  
Lasso Regression=0.062109979229048484  
ElasticNet Regression=0.220680994144546  
Logistic Regression=0.8741416915744405  
Random Forest=0.8745527986151183  
Logistic Regression is Suitable for this Dataset=0.8745527986151183
```